

Natural Language Processing with Disaster Tweets

Alekya Rama & Namitha
Ramakrishnan



Problem Statement

Objective: To Classify tweets as disaster-related (1) or not (0)

Goal:

- Use AI & NLP to distinguish real crisis events from casual tweets.
- Helps emergency teams filter urgent information faster.

Disaster - Related Tweet



Non Disaster- Related Tweet




Dataset Overview

Source: Kaggle's NLP Disaster Tweets dataset

Train Data: **7,613** labelled tweets

- Disaster (1): 43%
- Non -Disaster (0: 57%)

Test Data: **3,263** unlabelled tweets (for evaluation)



Key Data Challenges

- Noisy Text:
Misspellings, slang, emojis, and hashtags
- Short Texts:
280-character limit reduces context
- Figurative Language:
Hard to distinguish real disasters from expressions
- Missing Data:
Some tweets lack location or keyword

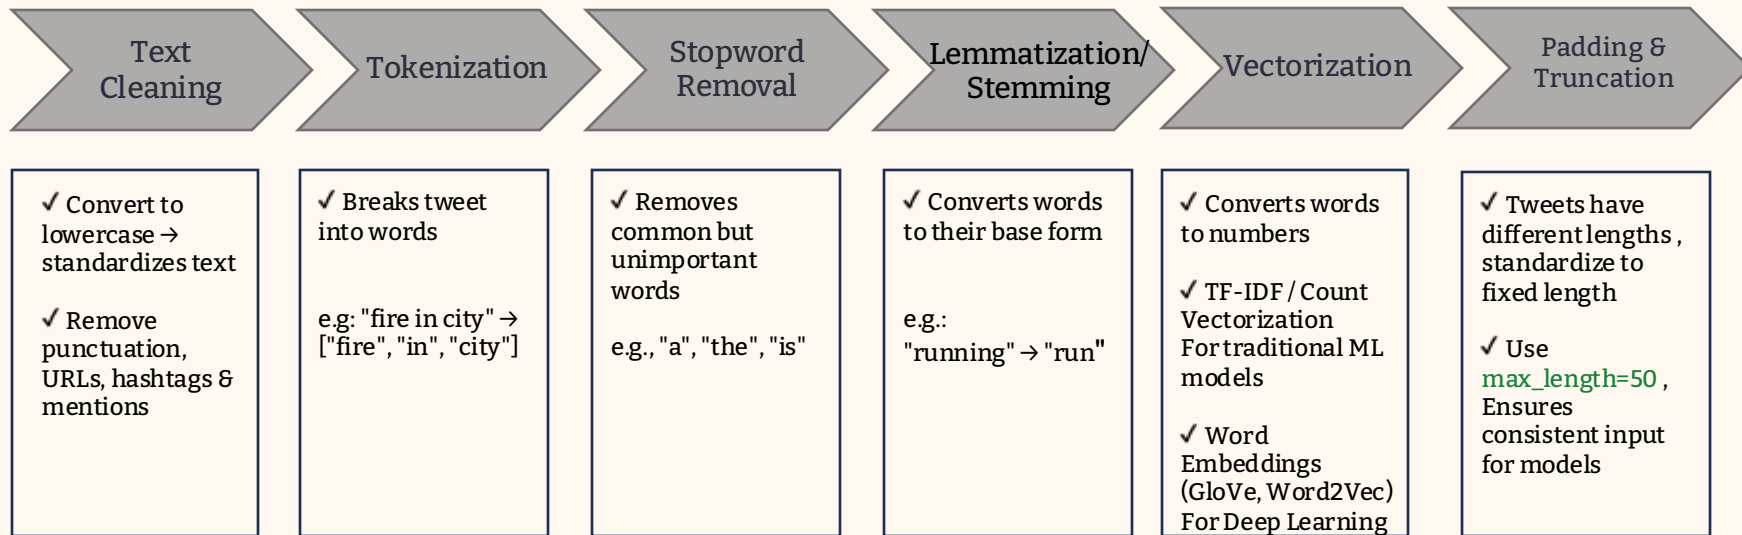
Data Pre Processing

Why Preprocessing Matters?

Raw tweets contain noise such as:

- Special characters & punctuation "Help!!! Flooding in Texas 😱😱"
- URLs & mentions "Check this out! <https://xyz.com> @user"
- Stopwords "is", "the", "at", "on" (common but non-informative words)
- Inconsistent casing "FIRE" vs. "fire"

Data Pre Processing



Model Performance Evolution

Model No.	Model Type	Accuracy	Key Enhancements
1	Baseline CNN/LSTM	0.770	<ul style="list-style-type: none">✓ Standard Text Pre processing✓ Simple architecture
2	CNN + LSTM	0.773	<ul style="list-style-type: none">✓ Added dropout✓ Batch normalization✓ Improved embeddings
3	Ensembler Model – CNN + LSTM	0.81	<ul style="list-style-type: none">✓ Increased LSTM Units,✓ Increased Regularization✓ Improved Learning Rate

Model 1 – BiDirectional LSTM Model (Accuracy: 0.770)

Features:


- Forward and Backward Processing
- Embedding Layer
- Dropout & Recurrent Dropout
- Stacked LSTM Layers
- Dense Output Layer (Sigmoid Activation)

Training & Validation Analysis:

- Good training accuracy, but validation accuracy plateaued
- Overfitting observed, as validation loss increased after a few epochs

Key Takeaways

- More dropout, L2 regularization, or even batch normalization layers can help reduce overfitting.
- If overfitting persists, consider reducing the model's capacity (e.g., fewer neurons or layers) or further tuning hyperparameters.
- Always monitor validation accuracy and loss. Relying on training metrics alone can be misleading.



Model 2 – Improved CNN + LSTM (Accuracy: 0.773)

Features:

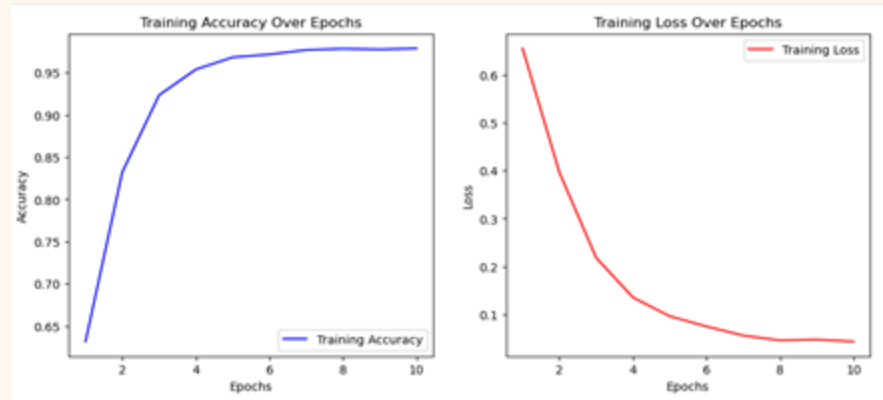
- Convolutional layers extract local patterns
- LSTM layer captures long-range dependencies in text
- Dropout (0.5) & L2 Regularization to control overfitting
- Optimized learning rate and batch size

Training & Validation Analysis:

- Slightly better generalization than previous model
- Reduced overfitting, but validation accuracy still fluctuated

Key Takeaways

- Better sequential understanding using LSTM.
- Still some overfitting—could be improved with better regularization



Model 3 – CNN + BiLSTM (Accuracy: 0.810)

Features:

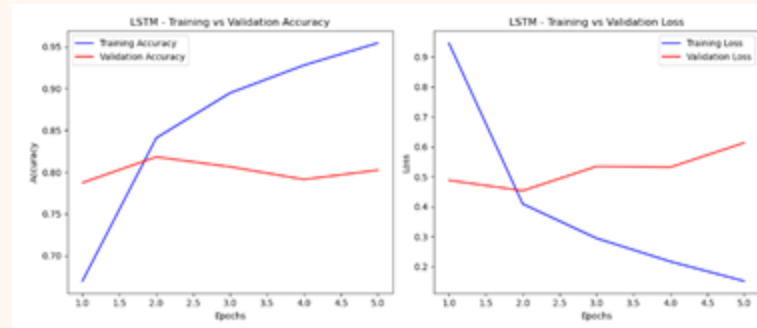
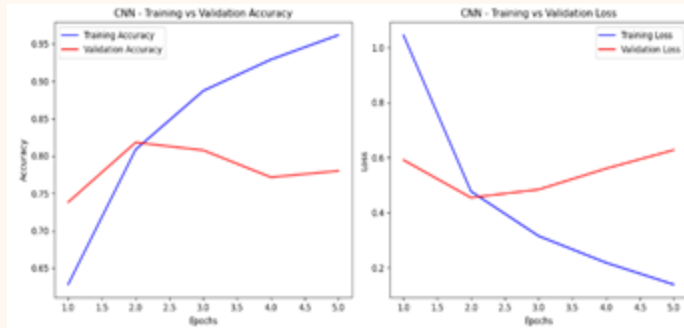
- Combined predictions from CNN & LSTM models
- Averaged final outputs instead of selecting a single model
- Early stopping & adaptive learning rate for optimization

Training & Validation Analysis:

- More stable validation accuracy, less overfitting
- Smoother loss curve compared to others
- Stronger robustness to tweet variations


Key Takeaways

- Best-performing model so far—ensemble approach improved generalization
- Still not 100% accuracy—potential for BERT or Transformer-based models



Key Takeaways:

Model Evolution:



Each step improved accuracy & robustness of the model

- **Overfitting Reduction** - Dropout layers, regularization techniques helped.
- **Generalization** - The ensemble model performed better on unseen data.

